

Master's Thesis: MuZero

Deep Reinforcement Learning with MuZero: Theoretical Foundations, Variants, and Implementation for a Collaborative Game

Content

Overview	3
AlphaZero	7
MuZero	26
Multiplayer Modifications	49
Evaluation	53
Questions	57


Overview

MuZero

- Model-based deep Reinforcement Learning algorithm
- Developed by Google DeepMind
- Games: Go, shogi, chess, Atari
- Evolution:
 - AlphaGo
 - AlphaGo Zero
 - AlphaZero
 - MuZero


AlphaZero and MuZero

- Learn from scratch using selfplay
-



AlphaZero and MuZero

- Learn from scratch using selfplay
- Use Monte Carlo Tree Search (MCTS) to plan ahead



AlphaZero


Neural Network

- Input: 2D image of game board
- Residual CNN
- Two output heads:
 - **value**: Scalar
How good is the current position?
 - **policy**: Distribution over actions
What are promising moves to try in the search?


Monte Carlo Tree Search

- Builds (game) tree of possible future actions
- Stochastic algorithm: Random samples in the action space
- Tree grows iteratively:
 1. **Selection:** *Find the most urgent node*
Guided by:
 - network policy predictions
 - exploration
 - exploitation
 2. **Expansion:** *Add a new node, query network*
 3. **Backpropagation:** *Update statistics in the tree*
- New search tree at every game state to find a move
- Search results are used as policy training target
→ **Policy improvement**


MCTS: Root Node




MCTS: Root Node: Network Inference




MCTS: Root Node




MCTS: Iteration 1: Selection




MCTS: Iteration 1: Expansion




MCTS: Iteration 1: Expansion (Network inference)




MCTS: Iteration 1: Backpropagation




MCTS: Iteration 1




MCTS: Iteration 2: Selection




MCTS: Iteration 2: Selection




MCTS: Iteration 2: Expansion




MCTS: Iteration 2: Expansion (Network inference)




MCTS: Iteration 2: Backpropagation



MCTS: Iteration 2



MCTS: After many Iterations



Selfplay and Neural Network Training

- Play games using MCTS for both players
- Record training data (s, π, z) for each game state
 - s : Game state Observation image
 - π : MCTS policy Distribution over actions
Which actions of the root node were visited by the search?
 - z : Game Outcome Scalar
Ended the game in a win, loss or draw?
- Supervised training on this data


MuZero

Summary

- Like AlphaZero, but no simulator in the tree search
- Instead: Learns a model of the environment dynamics
- Represents game states in a latent space

Neural Networks

- 3 Networks:




MCTS: Observation


s

x	o	
	x	
o		o


MCTS: Representation Network




MCTS: Root Node




MCTS: Root Node: Network Inference




MCTS: Root Node




MCTS: Iteration 1: Selection




MCTS: Iteration 1: Expansion




MCTS: Iteration 1: Expansion (Network inference)




MCTS: Iteration 1: Backpropagation




MCTS: Iteration 1




MCTS: Iteration 2: Selection




MCTS: Iteration 2: Selection




MCTS: Iteration 2: Expansion




MCTS: Iteration 2: Expansion (Network inference)




MCTS: Iteration 2: Backpropagation



MCTS: Iteration 2



MCTS: After many Iterations




Selfplay


- Play games using MCTS for both players
- Record training data (s, a, r, π, G) for each game state
 - s : Game state Observation image
 - a : Action taken Onehot tensor
 - r : Reward experienced Scalar
 - π : MCTS policy Distribution over actions
 - G : n-step return Scalar

Neural Network Training

- Sample K consecutive training steps from buffer
- Start with observation
- Then unroll dynamics network for $K - 1$ steps using actions
- Backpropagation-through-time
- End-to-end learning of policies, values and rewards K steps ahead



Neural Network Training (ii)



Multiplayer Modifications

Goals / Improvements over MuZero

- Multiplayer support
- Arbitrary turn order
- General-sum games
- Chance events / Stochasticity

Multiplayer MCTS

- At each node: Maximize current player's profit
- Requires:
 - Current player at turn
 - Predicted by the dynamics network
 - Trained on ground-truth labels from the game simulator
 - Per-player values and rewards
 - $v, r \in \mathbb{R}^n$ for n players
 - Trained on ground-truth labels from the game simulator


Stochasticity

- Add special chance player to the set of players
- Is at turn when chance events occur
- Policy targets are the chance outcome probabilities (ground-truth from simulator)
- MCTS selects actions according to predicted chance policy when the chance player is at turn

Evaluation


Carchess

- Cars spawn and drive along roads
- Cars may crash at intersections
- Players control traffic with barriers
- Goals:
 - Bring cars to their destination
 - Avoid crashes
- 10 Rounds of:
 - Each player toggles one barrier
 - Traffic advances for 5 steps
 - Spawn counts are updated randomly
- Collaboration: All players receive same reward



Carchess: Results

- Single training run
- 100.000 environment steps of random play at the beginning
- Mean score + 95% CI over last 30 plotted data points



Thanks

Questions